FROM WEB · DOC_V2_FROM_WEB

Conversation Summary & Document Index (V2)

VERSION
0.1 DRAFT

STATUS
DRAFT

OWNER
Hawkeye Platform

RENDERED
2026-06-11

CONTENTS

Generated: 11 June 2026 **For:** Local filing at `C:\Users\debab\Code - Hawkeye\hawkeye-clean\Doc_V2\from Web` **By:** Claude.ai chat sessions (curated)

---

1. What this bundle is

This zip contains the **current, canonical** set of Hawkeye strategy, product, and financial documents produced across multiple Claude.ai chat sessions ending June 11, 2026. It is **not** every draft ever produced — many earlier iterations were superseded by later versions, and only the latest is included here.

Total: 46 files, 19MB, organized in 7 numbered folders.

The bundle is designed to drop straight into your `Doc_V2/from Web/` directory and reflects the structure recommended for that location. The numbered prefixes are intentional — they're the reading order.

---

2. How to use this folder

IF YOU NEED TO...	OPEN THIS FIRST
Understand everything in one document	`01_canon/MASTER-REFERENCE.pdf`
Pitch an angel investor (deck + email)	`02_strategy_and_pitch/pitch/pitch-deck-angel-round.pdf` + `02_.../memo/founder-memo-2pp.pdf`
Run a customer meeting (product brochure)	`02_strategy_and_pitch/brochures/management-product-brochure-v2-LATEST.pdf`
Run a technical/IT meeting	`02_strategy_and_pitch/brochures/technical-user-deck.pdf`
Defend unit economics to a CFO	`05_financials/cost-model.xlsx` + companion PDF
Show the financial plan + cap table	`05_financials/business-and-funding-plan.pdf`
Demonstrate the architecture (verified)	`04_specs_and_research/pillars-architecture-VERIFIED.pdf`
Clean up your local docs folder	`07_cleanup_tools/claude-code-prompt-for-local-cleanup.md`

---

3. The conversation in one page

The work spans many sessions but lands on a coherent thesis. Here it is.

Strategic reframe (this matters more than anything else)

Hawkeye is NOT a pharma EQMS that might eventually expand to other industries.

Hawkeye IS an **industry-agnostic, AI-native compliance engine for regulated supply chains** — and pharma is the *forcing function*, not the destination. Pharma produces the hardest regulatory bar (GxP, 21 CFR Part 11, ALCOA+, validated AI, on-prem sovereignty), which forces an architecture every other regulated supply chain (food, devices, automotive, aerospace, chemicals) can inherit by configuration.

Seven elements regulated rigor forces structurally:

1. ALCOA+ data integrity
2. 21 CFR Part 11 e-signatures
3. Append-only audit trail
4. Tamper-evident hashing (SHA-256 — NOT blockchain)
5. Validated AI + per-call reproducibility
6. Model sovereignty / on-prem capability (validation-gated)
7. Separation of duties + RBAC + validation posture

What's actually built (code-verified)

- **170 Mongoose models** in `codex_backend_01` (public GitHub: `github.com/hawkeyetransact-dotcom/codex_backend_01`)
- **Two runtimes:** Python data-platform (crawlers, parsers, PDF extraction, entity resolution) FEEDS Node application
- **15 default modules:** AUDIT_MGMT, DOC_CONTROL, CAPA, CHANGE, EVENT, TRAINING, RISK, SUPPLIER_QUALITY, MGMT_REVIEW, ASSET, CHAIN_OF_CUSTODY, TRANSACTION_REVIEW, REGULATORY_INTEL, AI_ASSISTANT, RFQ_PROCUREMENT
- **Five-pillar architecture:** Collect → Process → Validate → Generate Report → Seal Record. Pillar 3 is `compliance/standardRegistryService` + `complianceEvaluationService` + `complianceRules.js`. Pillar 5 is `auditTrailService.js` + `buildSnapshotHash` (SHA-256).

- **Configuration layer:** `vocabularyService`, `universalModuleConfigService` (industryProfile='PHARMA_GMP'), `defaultStandards.js`, `modulePacks.js`, `WorkflowDefinitionService.js` — real, not vaporware
- **AI:** multi-LLM gateway (Anthropic/OpenAI/Gemini/local-vLLM), grounded generation, redaction, governance, 6 audit-agents, wave 2 (multi-step agent, custom tool-calling runtime — NOT MCP), wave 3 (on-prem LLM deploy with frozen `ONPREM_VALIDATION_REQUIREMENTS` IQ/OQ/PQ passThreshold 0.95)
- **SupplierRiskSnapshot:** 12-dimensional, fused continuously from public + internal data
- **Cross-module wiring:** `crossModuleService.js` makes deviations auto-create CAPAs, changes auto-assign training, audit findings flow into supplier scorecards

Honest gaps (the things NOT yet true)

- **Pre-customer.** No paying pharma logo references yet. The platform is real; the commercial proof is being built.
- **"Immutable" = tamper-evident** (SHA-256 + append-only audit trail), NOT blockchain. Important distinction for any technical conversation.
- **Live shared-audit marketplace** is a THIN module — framework built, network liquidity is the roadmap bet.
- **On-prem LLM deployment** has real scaffold + validation requirements but is NOT yet proven end-to-end in production.
- **SOC 2** is in flight (target Q4 2026); not certified today.
- **Some modules are partial:** Training, complaint, management review are functional but less mature than the core EQMS modules. RFQ-procurement is THIN.

---

4. The fundraising thesis (current)

The ask

- **Pre-seed / angel round: $1.2-1.5M** (revised down from the original $3M after building the proper business plan)
- **Use of funds (18 months):** 60-70% team build (10-15 FTE India), 10% AI infrastructure (~$115K for fine-tune compute + self-host), balance for cloud, compliance, marketing, ops
- **Path:** angel → seed at month 18 ($3-5M) → Series A at month 30-36 ($10-15M)
- **Cap table after Series A:** founders together hold ~47% (23.5% each from 50/50 start)

Market

- **India pharma manufacturing units:** ~10,500 total, ~3,000 WHO-GMP certified
- **Reachable beachhead:** 200-300 Tier 2 mid-pharma + Tier 3 CDMO + WHO-GMP-certified accounts
- **3-year SOM India pharma alone:** $1.4-2.3M ARR ceiling at blended $9.5K ACV
- **Ring 1 expansion:** Food & Beverage (architecturally near-identical), cosmetics, devices, blood/tissue
- **Global ceiling:** $3-5B SAM (regulated mfg at our ACV tier)

Unit economics (the corrected version)

The cost model showed prior documents understated cost-to-serve. The honest numbers:

CUSTOMER SIZE	VARIABLE COST/YR	FULL PRICE ACV	GROSS MARGIN
Small (1 site, 5 users)	$1,011	$4,500	77.5%
Medium (3 sites, 20 users)	$3,840	$10,800	64.4%
Large (5 sites, 50 users)	$11,543	$22,000	47.5%

Where the cost goes: Support (47% of variable cost) + AI inference (27%) dominate. Cloud infra is < 10%. Self-hosted AI in Year 2+ is the single biggest cost-reduction lever (~40% off variable cost).

Blended gross margin at scale (100+ customers): ~60% — *not the "~80%" cited in earlier documents*. This needs updating across the pitch deck, business plan, and management deck.

The Founding Customer offer

- **Year 1:** Software FREE; customer pays only implementation/config (~₹3-5L one-time)
- **Year 2+:** 50% of full software subscription + cost-of-serve passthrough (cloud + AI usage)
- **Sandbox:** lifetime access, separate from production, survives subscription end
- **Limit:** First 10 customers only, time-limited framing — so the discount is a program, not market price

Reality check: Year 1 is a LOSS of $1,311 (S) / $5,840 (M) / $17,043 (L) per customer once cost-of-serve is included. Frame it as "we absorb cost-to-serve as part of the program." Year 2+ rebuilds to 49-69% gross margin; payback 7-19 months by size.

---

5. File-by-file inventory

`01_canon/`

MASTER-REFERENCE.pdf (30 pp, 2.7MB) — *The definitive document.* Nine parts: vision, industry-agnostic engine reframe, regulated-rigor architecture, pharma domain depth, market reckoning, pharma plug-in spec, GTM, URS, honesty register. If anyone asks for "everything you have," send this.

`02_strategy_and_pitch/pitch/`

pitch-deck-angel-round.pdf (12 slides) — Angel-round pitch deck. ⚠ **NEEDS UPDATING:** currently asks $3M; the business plan now revises this to $1.5M. Same fix applies to unit economics claims ("80% margin" should be "60-78% by size").

`02_strategy_and_pitch/memo/`

founder-memo-2pp.pdf (2 pp) — Email companion to the pitch deck. Founder-voice. Reads like a real person wrote it, not a deck.

`02_strategy_and_pitch/brochures/`

- **management-product-brochure-v2-LATEST.pdf** (9 pp, 1.1MB) — *Latest, post-feedback rebuild.* Mobile-readable typography, six custom-generated infographics, quantitative value props with dual sourcing (industry benchmark + Hawkeye-modeled), NO commercials. Use this for management/exec conversations.
- **technical-user-deck.pdf** (11 slides) — Technical deck for QA/IT/CSV teams. No pricing. Pure product/architecture.
- **_archive_management-deck-v1.pdf** — Previous management deck. Kept because it has the Founding Customer offer details now removed from v2; useful as reference. Don't send.

`03_proposals/`

Currently empty — the POC and Commercial proposals were drafted in conversation but not all rendered to files in time. See "Pending work" below.

`04_specs_and_research/`

- **pillars-architecture-VERIFIED.pdf** (8 pp, 1.3MB) — Verified-against-code architecture document. The most defensible single thing we have for a technical due-diligence conversation.
- **URS-v1.0-DRAFT.pdf** (8 pp) — User Requirements Specification, Part A (foundational) + Part B (white-space).
- **audit-management-module-spec.pdf** (10 pp) — Detailed module spec for audit management.
- **quality-software-research-paper.pdf + .md** (4 pp) — Vendor-neutral industry research paper. Markdown source included for editing.
- **eqms-supplier-audit-industry-study.pdf** (8 pp) — Industry study with current SaaS landscape and six white spaces identified.

- **sector-market-analysis.pdf** (8 pp) — Per-sector market analysis with rings diagram and expansion sequence. Includes the "horizontal trap" honest reckoning.

`05_financials/`

- **cost-model.xlsx** — 6 sheets, 78 inputs, 300 formulas, zero errors. The unit-economics model. Edit blue cells; everything else recalculates. *Open `0_README` sheet first.*
- **cost-model-companion.pdf** (7 pp) — Narrative explanation of the model. The page-7 "honesty implications" table is where the corrections to prior documents are documented.
- **business-and-funding-plan.pdf** (11 pp) — Bottom-up market sizing, team build (Indian salary bands cited), native AI strategy (recommends fine-tune open-source over train-from-scratch), ROI-based pricing, 36-month financial model, two-round dilution to Series A.

`06_diagrams/`

The full visual canon. PNG + SVG for each.

- **pillars-asbuilt** — verified five-pillar architecture diagram
- **sector-rings** — pharma core to outer rings (food, devices, automotive, aerospace)
- **module-pillars-audit/capa/change/deviation/doc/supplier** — five-pillar walkthrough per module
- **audit-lifecycle, audit-admin-map, audit-data-model** — audit module specifics
- **v2_management_visuals/** — six new infographics generated for the v2 management brochure: audit-time bar chart, CAPA cycle comparison, time-allocation donuts, digital-thread flow, supplier-risk radial, four-numbers value summary

`07_cleanup_tools/`

- **hawkeye-chat-artifacts.zip** (6.9MB) — The earlier-built zip with 33 canonical files + MANIFEST.json, for merging into `backend/docs/` via the Claude Code prompt below
- **claude-code-prompt-for-local-cleanup.md** (166 lines) — Instructions for Claude Code to clean up your local `backend/docs/` folder hash-aware, with dry-run + safety rails. Use this if you want to consolidate older docs in `backend/docs/`.

---

6. Important corrections to prior documents

Several earlier-built documents contain claims that the cost model later showed were wrong or overstated. **Before sending any of these externally, fix:**

Pitch deck

- ❌ "~$1,900 variable cost / customer / yr" → ✅ "$1-12K depending on size"
- ❌ "~80% gross margin" → ✅ "60-78% gross margin depending on customer size"
- ❌ "Raising $3M" → ✅ "Raising $1.5M (per business plan derivation)"

Management deck (v1, archived)

- ❌ Founding Customer Year 1 framed as "implementation profit" → ✅ "We absorb cost-to-serve as part of the program; first 10 customers only" (the v2 brochure has this right)
- ❌ "4.5× ROI in 3-4 months" applies to *customer's* ROI on their savings, not the platform's margin — clarify which when speaking

Business plan

- ❌ The "~$1,900 / 80%" reference numbers — the bottom-up TAM and team build are fine, but unit economics should be updated to the cost-model numbers

---

7. What's still pending (not in this bundle)

These were discussed and partially drafted in conversation but not rendered to files by the time of bundling:

1. **Executive POC Proposal v2** — rebuild in the new mobile-readable style. The v1 content was good; needs typography refresh.
2. **User POC Proposal v1** — entirely pending. The "deep version" with full feature list, per-process and per-role depth, full market matrix.
3. **Commercial Proposal v1** — entirely pending. ROI math from cost model, Founding Customer offer detail, payback, TCO comparison.
4. **Updated pitch deck** with corrected $1.5M ask and corrected margin language.
5. **Updated business plan** with corrected unit economics.
6. **Filled-in founder credentials.** Every doc has `[ Founder name ]` placeholders. *Most important thing to fix before sending anything externally.* Generic team slides kill pre-seed deals.

---

8. The honesty discipline (preserved throughout)

This phrase keeps coming up in the documents — it's the through-line:

- Tamper-evident, **NOT** blockchain
- Custom tool runtime, **NOT** MCP
- On-prem LLM "available, validation-gated, **NOT** proven e2e"
- Industry-agnostic *by design*, pharma proven, other verticals = config roadmap
- AI presence is table stakes — the edge is the *combination*
- Live shared-audit marketplace = roadmap bet, **NOT** built network
- Pre-customer status acknowledged in every document

This isn't humility for its own sake. It's that investors do diligence, customers run PoCs, and inspectors verify. Naming what's real vs roadmap *first* is what makes the rest of what's claimed defensible.

---

9. How this was produced

This documentation was generated across roughly 6-8 hours of work in Claude.ai chat sessions between May 7 and June 11, 2026. The codebase verification used the public repositories `codex_backend_01` and `codex_frontend_01` cloned from `github.com/hawkeyetransact-dotcom/`. Industry benchmarks were web-searched in real-time and cited where used. Financial assumptions in the cost model come from public AWS Mumbai pricing (April 2026), published Anthropic/OpenAI/Gemini rates, 2026 India SaaS engineering salary benchmarks (Scaler, Glassdoor, ClarUP), and standard SaaS industry support-load assumptions.

Every number in this bundle is auditable. Every claim is sourced. Every limitation is named.

End of summary. Read the master reference document next, then the cost model, then everything else as needed.