

HAWKEYE

The compliance engine for regulated supply chains.

We collect supplier data, validate it against the rules, generate the audit report, and seal it as an inspection-ready record — in hours, not weeks. Built pharma-first; designed to work for any regulated industry.

Pre-seed / Angel round • April 2026

HAWKEYE

01 / 12

THE PROBLEM

A pharma contract manufacturer hosts ~30 audits a year from buyers and regulators — with a quality team of 5.

3-5 days PER AUDIT To prepare a single audit: gather evidence, build questionnaires, draft reports. Manual, repetitive, no reuse across audits.	100+ clients AUDITING THE SAME SUPPLIER Every buyer audits independently. Every supplier hosts the same audit dozens of times a year. Pure redundancy.	\$15B+ SPENT ANNUALLY On supplier quality + audit software globally. Growing 10-13%. Nobody owns the seam where the pain actually is.
--	---	--

The pain repeats across every regulated industry — pharma, food, devices, automotive, aerospace. Same shape, different standards.

WHY NOW

Three things are true in 2026 that weren't true two years ago.

- 01 AI is good enough to draft regulated documents** — and the governance models (grounding, citations, ISO 42001) now exist to make it defensible to an inspector. Before this, AI in GxP was unsellable.
- 02 The standards themselves are converging.** ISO 9001:2026, aerospace IA9100, automotive IATF 16949 are all re-aligning to a common spine. The regulatory world is moving toward exactly the engine-plus-config architecture we built.
- 03 The cross-industry seam is unowned.** Pharma EQMS vendors (Veeva, MasterControl) are deep but pharma-specific. Audit networks (Qualifyze) are network-only with no EQMS. Nobody runs both, across industries.

THE INSIGHT

Regulated compliance is the *same five-step motion* in every industry:

collect data, validate against a standard, generate the report, seal the record.

What changes is the *standard* — not the engine.

Build the engine once, to the hardest regulatory bar (pharma GxP), and every other industry becomes a configuration pack. That's a 10x category, not a feature.


WHAT WE BUILT

The full pharma plug-in is shipped. 170 data models, 15 modules, 20+ AI agents.

INTERNAL QUALITY Document control, deviation, CAPA, change, training, complaint, risk, management review — closed-loop, e-signed, audit-trailed.	AI THAT AN INSPECTOR ACCEPTS Every AI output cites its source, records the prompt + model + confidence, and is signed by a human. Per-call reproducibility — not a black box.
SUPPLIER AUDIT Risk-weighted questionnaires, AI-drafted dossiers, public-data fusion (FDA/EMA/WHO), evidence capture, cited reports.	21 CFR PART 11 + ALCOA+ SHA-256 tamper-evident records, append-only audit trail, separation of duties — built from day one.
THE THREAD BETWEEN A deviation auto-spawns a CAPA. A risk event triggers a for-cause audit. An audit finding flows into supplier scorecard. Nobody else couples these.	SOVEREIGNTY OPTIONAL Cloud, hybrid, or on-prem LLM — for buyers who can't let process data leave their boundary.

HOW IT WORKS

One engine, five pillars, pluggable AI at every step.



VERDICT: all 5 pillars exist as discrete, real components. Pillars 1-2 = Python data-platform; 3-5 = Node app. Standards & vocabulary are DATA (config), confirming engine/config separation.
Honest nounce: AI attaches per-module & per-wave (not literally one slot per pillar); validation today is audit/compliance-centric; on-prem is validation-gated not yet proven a2c; toolCalling/runtime is custom (not MCP).

Configuration layer (top) makes the same engine pharma today, food tomorrow, automotive after that. AI plugs into each pillar — swap models, run on-prem, never sign for a human.


WHY THIS IS HARD TO COPY

Not any single feature — the **combination** nobody else has.

1. EQMS + AUDIT-NETWORK COUPLING Pharma EQMS vendors don't run audit networks. Audit-network players don't run EQMS. Building either side from scratch takes 5+ years.	2. PER-CALL REPRODUCIBLE AI Every output traces to its prompt, model, sources. Most "AI EQMS" is opaque — ours is built for the inspector.
3. LIVE PUBLIC-DATA FUSION FDA, EMA, WHO signals stream into supplier risk continuously. Most platforms use self-reported, point-in-time data.	4. INDUSTRY-AGNOSTIC ENGINE Standards as data, not code. Adding food / automotive is configuration, not a rebuild. Vertical incumbents structurally can't do this.

Each piece exists somewhere. **The combination doesn't.** That's the moat.

Pharma first — then ring by ring outward.



BUSINESS MODEL

Usage-aligned SaaS. No big upfront. Customers pay as they get value.

Free 2 PROOFS-OF-CONCEPT Customers try us free on real data for 3-4 weeks before paying. Conversion is the goal.	Per site + per user SUBSCRIPTION BASE Annual platform fee per facility + named-user fee. Read-only viewers are free.	AI credits VARIABLE Monthly allowance + overage for additional AI usage. Credits weighted by action complexity.
~\$9,500 TARGET ACV / CUSTOMER Sized to fit India's RPTUAS subsidy. ~⅓ the price of Veeva/MasterControl. Deliberate.	~\$1,900 VARIABLE COST / CUST / YR At steady state. Cloud infra (~\$1,620) + AI inference (~\$280).	~80% GROSS MARGIN Healthy unit economics from day one. Scale collapses cost-per-customer further.

COMPETITION

We're the only column with **no gaps** — spanning where each specialist can't.

Capability	Veeva / MasterControl	CQ / Qualio	Qualifyze / Rx-360	India tier (AmpleLogic)	Hawkeye
Full EQMS depth	✓	✓	✗	✓	✓
Live supplier-audit network	✗	✗	✓	✗	~
Per-call reproducible AI	~	~	~	✗	✓
Public-data risk fusion	~	✗	~	✗	✓
EQMS↔audit thread	~	~	✗	✗	✓
Industry-agnostic engine	✗	✗	✗	✗	✓
Affordable for SMB / emerging markets	✗	~	~	✓	✓

We're not trying to beat Veeva at being Veeva. We're the only player spanning all of it — including the cross-industry primitive nobody else can build without abandoning their moat.

TEAM

The people building this.

FOUNDER & CEO

[Your name]

[Background — relevant industry / product / engineering experience. Why *you* are the right person to build this. Specifics: years in pharma, prior companies, technical depth.]

CO-FOUNDER & CTO

[Co-founder name]

[Background — technical depth, prior systems built, AI/ML credentials if relevant. The architect behind the verified five-pillar engine.]

ADVISORS

[Named pharma quality leaders, regulatory advisors, GTM advisors — concrete names matter a lot here.]

ENGINEERING / PRODUCT

[Team size & composition. Engineers + domain experts. The code lives in production-grade repos with 170 models and a verified architecture.]

WHY US

[One paragraph: the founder-market fit. Why this team, this problem, this moment.]

⚠ **Founder note: this slide is a placeholder. Team is typically the #1 driver of a pre-seed decision — fill it with real names, specific credentials, and verifiable experience before sending. Generic team slides kill deals.**

Raising \$3M to take pharma proven and build the cross- industry edge.

USE OF FUNDS • 18-24 MONTHS

40%

Product & engineering — harden pharma pack; ship F&B config; close white-space items.

15%

Compliance & validation — CSV docs, security certs, on-prem validation packs.

35%

GTM — design partners → first 20 paying customers in pharma; reference deployments.

10%

Reserves / runway buffer.

M 6

5 paying pharma customers; first ref deployment live; SOC 2 underway.

M 12

20 paying customers; F&B pack shipped; first non-pharma design partner.

M 18

Cross-vertical proof point; ARR run-rate to justify Series A on real numbers.

Let's talk — founder@hawkeye.app