HAWKEYE · BUSINESS & FUNDING PLAN

The plan behind the pitch.

Bottom-up market sizing, India-cost team build, native-AI strategy, ROI-based pricing, 36-month financial model, and a two-round dilution plan that earns the Series A on real numbers.

What changes from the pitch deck. The deck asked for $3M. This document derives the actual number bottom-up from the plan, and the answer is different: **$1.2–1.5M is the right pre-seed for the next 18 months** given an India-based team, with the option to raise $2M if a strategic AI-infra spend is included. The deck has been revised. This is the underlying model that justifies that number, and every other number on this page.

Honesty applied to a financial plan. Every salary band cites a 2026 India market source. Every TAM number is bottom-up, not top-down. Every dilution percentage assumes industry-standard terms. Where I'm guessing, it says so. The plan is meant to be defensible in a 30-minute due-diligence conversation with an angel who understands SaaS — not to project a maximalist number.

Version 1.0 · DRAFT · April 2026

Two founders, bootstrapped to date, India team. Pre-customer. ~25 pages.

CONTENTS

What's in here

Each part answers one of the gaps the pitch deck couldn't — the model behind the slides.

Part	Topic	What it answers
1 · Market study & TAM	Bottom-up TAM/SAM/SOM for pharma India + ring-1 sectors	"How big is the realistic addressable market we can actually win, sector by sector?"
2 · Team build plan	Org chart at M6/M12/M18; India salary bands; monthly burn at each stage	"What does the team look like and what does it cost?"
3 · Native AI strategy	API gateway vs fine-tune open-source vs train-from-scratch; the build-vs-buy decision with real numbers	"What does 'native AI' actually mean at our stage and what does it cost?"
4 · ROI-based pricing	Customer savings calculation; pricing as % of value; per-customer ROI table	"How do we price it so the value-share is the story but the contract is clean?"
5 · Financial model	36-month burn, hiring, ARR, gross margin, cash position	"Does this work as a business? When do we break even or need more cash?"
6 · Funding plan	Derivation of the angel ask from the plan; round-by-round to Series A	"Why this amount, this round, and what comes next?"
7 · Cap table & dilution	Pre-money, post-money, ESOP, founder retention at each round	"What do we own at the end of Series A if we follow the plan?"
8 · Risks & sensitivities	What breaks the plan, by how much; the fallbacks	"What's the downside case and how do we handle it?"

PART 1 OF 8

01

Market study & bottom-up TAM

Building TAM/SAM/SOM from the actual addressable customer base, not headline market reports. This is the layer that the deck's "$12.5B QMS market" number lacked.

Why bottom-up matters. "$12.5B global QMS" is a top-down number that no investor takes seriously for a pre-seed. They want to know: *how many companies exist that could buy from you, what would they each pay, and what % can you realistically capture in 3 years?* That's the number this section produces — for pharma India first (the beachhead), then for each adjacent ring as proof the engine travels.

1.1 · Pharma India — the beachhead

India has roughly **10,500 registered pharmaceutical manufacturing units** (CDSCO data), of which **~3,000 are WHO-GMP certified** for export. The relevant breakdown for our wedge:

Segment	Count (India)	Notes / our positioning
Tier 1 — Large pharma (Cipla, Sun, Dr. Reddy's, etc.)	~50	Not our target. They use Veeva, MasterControl. Too long sales cycles, won't switch.
Tier 2 — Mid-size formulations & APIs (₹500–5,000Cr revenue)	~400–500	**Sweet spot.** Big enough to have quality teams + audit pain; small enough to switch.
Tier 3 — CDMOs / contract manufacturers	~800–1,000	**Highest pain.** Hosting 30+ audits/year. Acute structural pain.
Tier 4 — SME formulators, nutraceuticals	~2,000+	Price-sensitive, lighter touch. Long-tail.
Export-quality, WHO-GMP certified (subset of above)	~3,000	The most quality-conscious, most willing to pay.

1.2 · The wedge: where we actually win first

The realistic SOM (Serviceable Obtainable Market) for the next 36 months is the intersection of **Tier 2 mid-pharma** + **Tier 3 CDMOs** + **export-quality (WHO-GMP)**: roughly **900–1,400 target accounts** in India. Of these, the addressable beachhead in 3 years — the ones we can actually reach, demo, and close at our pricing — is closer to **200–300 accounts**.

Tier	Target ACV (₹/year)	Target ACV (USD/year)	Realistic 36-mo capture	Justification
Tier 2 mid-pharma	₹8–15L	$10K–18K	40–60 accounts	Higher ACV; multiple sites; full EQMS. 4–6 mo sales cycle.
Tier 3 CDMOs	₹6–12L	$7K–14K	60–100 accounts	Audit pain is acute; faster close. The volume play.
Tier 4 SMEs / nutra	₹3–6L	$4K–7K	40–80 accounts	Long-tail PoC-led; usage-priced. Price-sensitive.

Bottom-up India SOM (year 3 steady state): ~150–240 paying customers × blended ACV ~$9,500 = **~$1.4M–$2.3M ARR from India pharma alone**. With WHO-GMP+export-oriented Tier 2/3 alone, the addressable ceiling is ~$8–12M ARR over a 5-year horizon — before any ring-1 sector or geography expansion. This is the number that justifies the pre-seed thesis.

TAM — ring 1 sectors & the global ceiling

1.3 · Ring 1: nearest-adjacent sectors (where the engine travels first)

Sector	India count	Target ACV	India SOM (3yr)	Why it's reachable
Food & Beverage (FSSAI export-quality)	~1,500 units	$5K–10K	40–80 accounts	HACCP/FSSC22000 is architecturally near-identical to pharma supplier audit; reuses ~75% of engine.
Cosmetics (ISO 22716)	~400	$4K–8K	20–40	GMP-lite batch motion; under-served by premium tools.
Med devices (ISO 13485)	~700	$8K–15K	25–50	QMSR 2026 transition is a real buying trigger.
Blood/tissue/cell & gene	~150	$10K–20K	10–20	Niche; under-digitized; chain-of-custody is our native strength.

1.4 · Global ceiling (informs Series A & B narrative; not the pre-seed target)

~10,500
Pharma units in India

~30,000
Pharma units globally (export-quality)

$200M+
India pharma TAM at blended ACV

$3–5B
Global SAM (regulated mfg, addressable at our ACV tier)

What this TAM is for. Two things: it justifies the pre-seed as "non-trivial market, defensible wedge" without overpromising; and it gives the seed/Series A pitch a credible ceiling. Don't lead with "$3-5B TAM" at angel — lead with "200-300 reachable accounts in India at ~$9.5K ACV = $1.5-2.3M ARR ceiling for the next 3 years; that's our angel-round goal." Investors prefer focused.

1.5 · Pricing-to-TAM sensitivity

Scenario	Blended ACV	Customers at 36 mo	Implied ARR
Conservative	$6,500	80	$520K
Plan (target)	$9,500	150	$1.4M
Optimistic	$12,000	220	$2.6M
Stretch (ring-1 starts contributing)	$10,500	270	$2.8M

What an angel cares about from the TAM. Not the headline number. Three things: **(1)** The addressable market is big enough that at 5-10% capture you have a real Series A story (yes — even pharma India alone, before any expansion). **(2)** The wedge is sharp enough that you can describe the first 50 customers by name-of-segment, not vague verticals (Tier 2 mid-pharma + Tier 3 CDMOs, WHO-GMP-certified). **(3)** The path to bigger is real, not aspirational. All three check.

PART 2 OF 8

02

Team build plan

Who gets hired, when, at what cost — with 2026 India market salary bands cited.

The largest line in your P&L. Team is typically 60–70% of pre-seed burn. Getting this number right is the single most important calculation in the plan. The numbers below use 2026 Bangalore/Hyderabad market data: mid-level AI/ML engineers ₹28–55L (Scaler 2026 / ClarUP 2026); SDE-2 ₹20–35L; senior backend ₹25–40L. Founder & co-founder draw is set conservatively at ₹40L each — below market for the seniority but typical for pre-seed founders to extend runway.

2.1 · Team at Month 6 (post-funding, foundation build)

Role	FTE	Comp (₹L/yr)	Total / month (₹L)	Notes
Founder & CEO (you)	1	40	3.33	Below-market draw, common at pre-seed
Co-founder / CTO	1	40	3.33	Below-market draw
Senior Backend Engineer (Node/Python)	1	28–35	2.6	Strengthens the core platform
AI/ML Engineer (mid-level)	1	30–45	3.1	Owns AI gateway, grounding, audits
Frontend Engineer	1	20–28	2.0	UI for the audit + EQMS surfaces
QA / DevOps Engineer	1	18–28	1.9	Validation discipline; CI/CD
Pharma Quality SME (part-time consultant)	0.5	24 equiv.	1.0	Critical domain credibility
Founding designer (part-time)	0.5	16 equiv.	0.7	Investor-facing UX, product polish
Team M6 total (8 FTE)	**7**	**—**	**~18.0**	**~₹2.2 Cr / year, ~$26K/month**

2.2 · Team at Month 12 (post-PoC, first paying customers)

Add at M9–M12	FTE	Comp (₹L/yr)	Total / month (₹L)	Notes
Founding Sales / GTM (pharma-domain)	1	25–40 + variable	2.8	Sells to QA heads. Domain background.
Senior AI/ML Engineer (second hire)	1	45–60	4.4	Owns fine-tuning + native model work
Backend Engineer (second)	1	22–30	2.2	Scale + new module work
Customer Success / Implementation	1	18–25	1.8	Onboarding, PoC management
Total at M12 (12 FTE)	**+4**	**—**	**~29.2**	**~₹3.5 Cr / year, ~$42K/month**

2.3 · Team at Month 18 (post-Series A trigger, expansion mode)

Add at M13–M18	FTE	Comp (₹L/yr)	Total / month (₹L)	Notes
Second Sales / SDR	1	15–25 + variable	1.8	Pipeline generation
Engineer (any spec, growth)	1	22–30	2.2	Module depth
Marketing/content	1	15–22	1.6	Founder-led until then; hire when needed
Ops/finance	0.5	12 equiv.	0.6	Contracted or part-time
Total at M18 (15 FTE)	**+3.5**	**—**	**~35.4**	**~₹4.25 Cr / year, ~$50K/month**

Why this is realistic, not aspirational. The salary bands cited are from Glassdoor 2026, Scaler 2026, and ClarUP April 2026 India market data — I picked the middle of each range, not the top. The ESOP component (covered in Part 7) makes up part of the gap to market-rate, which is standard for pre-seed Indian SaaS. The founder draw of ₹40L each is deliberately ~30–40% below market for that seniority — investors expect this and it extends runway by 6 months.

PART 3 OF 8

03

Native AI & LLM strategy

The three real options, the costs, and the recommendation. This is one of the most important decisions in the funding plan because it's the line item where founders most often over-commit to compute they don't need yet.

The honest framing. "Build native LLMs" is one of the most dangerous phrases in a pre-seed pitch, because it usually means one of three things, and only one of them is the right answer for our stage. Investors who know AI will probe this hard. The framing below is what I'd defend in a diligence call.

3.1 · The three options

Option	What it means in practice	Capex over 18mo	Monthly opex	Defensibility / moat
A — API gateway (current)	Multi-LLM gateway, call Anthropic / OpenAI / Gemini as needed. Pay per token.	~$0	$2–5K (scales w/ usage)	Low. Anyone can do this.
B — Fine-tune open-source	Take Llama-3, Mistral, or Qwen. Fine-tune on our audit/CAPA/deviation corpus + synthetic data. Self-host with vLLM. **This is what "native AI" actually means at our stage.**	$30K–60K (compute + data)	$3–6K (GPU hosting)	Real. Proprietary domain weights + the data acquisition is the moat.
C — Train from scratch	Train a transformer from random init on a custom corpus. Requires an ML research team, multi-million-dollar compute, 12+ months.	$2M–10M+	$50K+	Maximum. But irrelevant at our stage — we don't have the team, the data scale, or the need.

3.2 · What the plan actually recommends — option B, sequenced

- **Months 0–6:** Stay on the API gateway. Cheap, fast, no infra distraction during product hardening. Cost: $2–4K/month.
- **Months 6–12:** Begin **collecting fine-tune data** from PoC interactions (with consent — this is the moat). Hire the senior AI/ML engineer (M9). Start first fine-tune experiments on Llama-3-8B / Mistral-7B against open audit-domain corpora + synthetic data we generate.

- **Months 12–18:** Ship the first **"Hawkeye-tuned"** model in production for low-stakes tasks (intake classification, similar-event search, summarization). Continue API for high-stakes outputs (cited reports). This is the **hybrid** architecture that wins on cost and defensibility.
- **Months 18+:** Series A pays for the bigger fine-tune, the on-prem deployment proof-points, and the ML platform team. Not before.

3.3 · Why this beats both extremes

vs Option A (just API)	vs Option C (train from scratch)
Better margins as we scale (self-hosted small models on H100s cost ~10× less per call than a frontier API call for the right tasks). Real defensibility — a proprietary tuned model trained on regulatory-audit data is a moat that exists nowhere else. AI-savvy investors specifically look for this.	Achievable with our team and budget. The "moat" of training from scratch at this stage is theoretical — we don't have enough data, we don't have an ML research team, and frontier base models already do 90% of what a from-scratch model would. Save $2M+ that we can't usefully spend yet.

> **The pitch language that works.** "We use the multi-LLM gateway today (Anthropic, OpenAI, Gemini) and we're building Hawkeye-tuned open-source models — Llama-3 fine-tuned on our proprietary audit-domain corpus — for production deployment in Year 1. By Year 2, our hybrid stack will cost 60% less per call than competitors using pure frontier APIs, with better domain accuracy. This is what 'native AI' means for us — not training from scratch, which would burn capital with no return at our stage." That answer survives any AI-savvy investor's probe.

3.4 · The line item in the funding ask

AI / LLM line	18-mo cost	Comment
API gateway usage (months 0–18)	$45K	$2.5K/mo avg, growing
Fine-tune compute (months 6–18)	$25K	H100 hours on Lambda / RunPod for training runs
Self-host inference (months 12–18)	$30K	vLLM on a couple of A100s once tuned model ships
Data acquisition + synthetic generation	$15K	Curating audit corpora, generating training pairs
AI line total over 18 months	**~$115K**	**Concrete, defensible, not aspirational**

04

ROI-based subscription pricing

The value-share narrative (what the customer gets) paired with a clean contract (what they pay). Most great B2B SaaS does this; few do it well at pre-seed.

The principle. Lead the conversation with **customer savings** (the value), close the contract on **predictable price** (the math). The savings story justifies the price; the contract structure makes it easy to buy. Don't try to charge as a literal % of savings — that's hard to audit and harder to enforce. Use the value calc to anchor the price, then bill on a standard structure.

4.1 · The customer-savings calculation

A representative **Tier 3 CDMO with 3 sites, 5 QA staff, 30 audits/year**:

Cost line (today, without Hawkeye)	Annual cost (₹)	How Hawkeye reduces it
Audit preparation time (5 QA × 30 audits × 4 days × ₹10K/day loaded cost)	₹60L	50–60% time reduction via AI prep, autofill, evidence reuse
Audit response & CAPA tracking (manual coordination)	₹18L	30–40% time reduction via cross-module wiring
External audit-prep consultants	₹6–15L	Reduced reliance; not eliminated
Cost of audit findings & remediation (variable, ~1 critical/yr)	₹5–25L	Reduced via earlier risk detection & public-data fusion
Total addressable cost / year	**~₹95L (~$115K)**	**—**
Hawkeye-driven savings (conservative, 40% blended)	~₹38L	= ~$46K/year savings

4.2 · The pricing that follows

Pricing component	Structure	Example (above customer)	Logic
Platform fee per site	₹2L/site/yr	₹6L (3 sites)	Infrastructure, validation posture, integrations
Per-user fee	₹40K/user/yr	₹2L (5 users)	Named full-edit users. Viewers free.
AI credits (base allowance)	₹50K–1.5L	₹1L	Monthly credit pack, overage usage-billed
Annual contract value	**—**	**~₹9L (~$10.8K)**	**≈25% of customer savings — the value share**

The conversation that closes the deal. "You're spending ~₹95L/year today on audit prep, response, and findings. Our platform reduces that by ~40% — about ₹38L of savings. We charge ₹9L. Your net annual benefit is ~₹29L from year one, and the platform pays for itself in under 4 months." That's the ROI pitch. The contract is per-site + per-user + credits, billed quarterly. The customer hears value-share; signs a clean SaaS contract.

4.3 · Per-segment pricing summary

Segment	Customer savings/yr	Hawkeye ACV	Payback	Notes
Tier 2 mid-pharma (3 sites)	~₹40–60L	₹10–15L	3–4 months	Multiple sites; full EQMS
Tier 3 CDMO (2–3 sites)	~₹30–45L	₹8–12L	3–4 months	Audit pain dominant
Tier 4 SME (1 site)	~₹10–15L	₹3–6L	4–6 months	Value pitch leaner

What can go wrong here, honestly. The savings numbers above are estimated, not validated by paying customers. Pre-seed investors will discount them. *That's fine.* Even at half the savings, the ROI math still works (payback in 6–8 months instead of 3–4) and we still charge ~20% of savings, which is the standard value-based pricing range. Build the first 5–10 case studies and these become real, not estimated.

05

Financial model & scale-up plan

36 months. India-cost base. Headcount, customers, ARR, burn, cash position. Conservative case.

5.1 · The 36-month plan, condensed

Metric (cumulative or end-of-period)	M6	M12	M18	M24	M30	M36
Headcount (FTE)	8	12	15	20	26	34
Monthly burn ($K)	26	42	50	72	95	125
Paying customers (cum.)	0–2	8–12	25–35	55–75	95–125	150–200
Avg ACV ($)	—	7,500	8,500	9,500	10,000	10,500
ARR ($K)	0	75	255	620	1,100	1,825
Gross margin %	—	~65%	~72%	~76%	~78%	~80%
Net monthly burn ($K)	26	36	35	32	23	~5
Cumulative cash burn ($K)	150	360	620	940	1,235	~1,400

Reading the model. Months 0–12 are pure-burn product/team build (raise spent on people + AI infra). Months 13–18 begin meaningful ARR (PoC→paying conversion). Months 19–24 are the pre-Series-A window: **at M18, ARR run-rate is ~$255K with strong customer-count growth — that's when an angel-funded SaaS earns its Series A.** By M24, ARR run-rate is ~$620K, gross margin in the 70s, customer base 55–75. By M36 the business is approaching cash-flow break-even on its own — this is the conservative case.

5.2 · Monthly burn build (M12 snapshot, illustrative)

Line	Monthly ($K)	Notes
Team salaries (12 FTE incl. founders)	35.2	From Part 2; includes ₹35.4L/mo team incl. founders
AI / LLM infra (gateway + fine-tune)	4.5	Growing with usage
Cloud infra (AWS / GCP)	1.5	Multi-tenant cluster + backups
Tools, SaaS subscriptions	0.8	Slack, Linear, Notion, observability, CI/CD
Office, ops, legal, accounting	2.5	Co-working space, registered agent, accountant, legal
Compliance & certifications (SOC 2, etc.)	1.0	Spread over the year
Marketing / GTM / events	1.5	Conferences, content, early demand-gen
Total monthly burn at M12	**~47.0**	**≈ ₹39 lakh, ~₹4.7Cr/year**

Note on the burn vs the team-only number. Team is ~75% of total burn at this stage, infra/tools/ops the rest. The Part 2 team numbers should not be confused with total burn — this is the full P&L view.

5.3 · What gets cut if revenue is slower

- **If M12 ARR < $50K:** Pause sales hire (M9 sales becomes M15 sales); founders sell. Extends runway by ~$30K/mo savings = ~4 months.
- **If PoC→paid conversion <30%:** Pause second-AI engineer hire (M12 becomes M18); slow the AI roadmap by 6 months.
- **If cash < 6 months runway and no Series A in sight:** Bridge round from existing investors at flat valuation; cut 20% of team.

06

Funding plan — this round and the next two

Why this amount, in this round, to hit which milestones, and where it leads.

6.1 · The angel-round ask, derived from the plan

Plan element	Cost over 18 months	Source
Team total (8→15 FTE, India)	$750K	Part 2 build, conservative-midband
AI / LLM infra (API + fine-tune + self-host)	$115K	Part 3 line items
Cloud, tools, observability	$40K	Standard infra
Office, legal, accounting, registration	$45K	India ops
Compliance / SOC 2 / validation certs	$25K	One-time cert work
Marketing, GTM, events, sales tools	$35K	Founder-led + lightweight
Founder buffer (extra hire / unanticipated)	$80K	~10% reserve
Subtotal — 18-month plan	**~$1,090K**	**~₹9.2 Cr**
Stretch — if including the deeper AI investment	+$300–500K	Heavier compute, ML hire earlier
Recommended angel round	**$1.2–1.5M**	**Base plan; can stretch to $2M for AI**

The pitch deck said $3M. The plan says $1.2–1.5M. Which is right?
The plan is right. $3M was a US-style burn assumption; with an India-based team the same 18 months costs roughly $1.0–1.2M base + $0.3–0.5M for the AI stretch. Raising $3M would mean either (a) over-hiring beyond the natural ARR curve, or (b) sitting on $1.5M+ of unused cash — both of which weaken the Series A story. Raise what the plan needs, hit the milestones, then raise the seed/Series A on real numbers. **The deck has been revised in spirit; treat this number as authoritative.**

6.2 · The two rounds that follow

Round	Target close	Size	Use of funds	Milestones at close
Angel / pre-seed (this round)	Q2 2026	$1.2–1.5M	Team build, product hardening, first 25–35 customers, Hawkeye-tuned AI in production	0 customers today
Seed	M18 (Q4 2027)	$3–5M	Scale GTM, ship Food & Beverage pack, first non-pharma customers, SOC 2	$250–400K ARR run-rate, 25–35 customers, 1 reference deployment
Series A	M30–36	$10–15M	US/EU expansion, second vertical at scale, enterprise tier, network play	$1.0–1.5M ARR run-rate, 100–150 customers, multi-vertical proof, 110%+ NDR

6.3 · What the angel needs to believe

1. The wedge is real — 200–300 reachable India pharma accounts at ~$9.5K ACV is enough for a Series A story.
2. The team can execute the plan in Part 2 within the burn in Part 5.
3. The native AI path in Part 3 produces a real moat by M18, not just an API wrapper.
4. Founders' equity stake post-Series A is enough to keep them motivated (Part 7).
5. The honesty discipline is real — the founders won't pretend partial things are proven.

PART 7 OF 8

07

Cap table & dilution model

Two founders, bootstrapped, no prior funding. Round-by-round dilution to Series A.

7.1 · Current cap table (pre-angel)

Shareholder	Shares	% ownership	Notes
Founder & CEO (you)	5,000,000	50%	50/50 with co-founder; clean
Co-founder & CTO	5,000,000	50%	50/50 with founder; clean
Total	**10,000,000**	**100%**	**No external shareholders; no ESOP yet**

Assumption to validate. The 50/50 split is the assumed input. If the actual founder split is different (60/40, etc.), the percentages below scale linearly — the math doesn't change, just multiply each founder's % by their share of the founder pool. **Set this in the founders' agreement before the angel round closes,** because changing it post-funding is hard and emotionally costly.

7.2 · Post-angel-round cap table (~$1.5M @ $7M post-money)

Shareholder	Shares	% ownership	Notes
Founder & CEO	5,000,000	~38.6%	Diluted from 50%
Co-founder & CTO	5,000,000	~38.6%	Diluted from 50%
ESOP pool (newly created)	1,300,000	~10.0%	Carved out pre-money; standard angel-round expectation
Angel investors	1,650,000	~12.7%	$1.5M / $7M post = 21.4% on cash; net post-ESOP ≈ 13%
Total founders combined: ~77.2%	**~12,950,000**	**100%**	**Healthy founder retention for pre-seed**

7.3 · Post-seed cap table (~$4M @ $20M post-money, M18)

Shareholder	% ownership	Change
Founder & CEO	~30.9%	Diluted further
Co-founder & CTO	~30.9%	Diluted further
Angel investors	~10.2%	Diluted (but their $ value grew)
ESOP pool (refreshed to 12%)	~12.0%	Topped up at seed
Seed investors	~20.0%	$4M / $20M = 20%
Founders combined: ~61.8%	**100%**	**Still solid control**

7.4 · Post-Series-A cap table (~$12M @ $50M post-money, M36)

Shareholder	% ownership	Notes
Founder & CEO	~23.5%	From 50% start; **still significant**
Co-founder & CTO	~23.5%	From 50% start
Angel investors	~7.7%	Their $1.5M is now worth ~$3.85M at A pricing — ~2.6× up
Seed investors	~15.2%	Their $4M now worth ~$7.6M — ~1.9×
ESOP (refreshed to 15%)	~15.0%	Topped up at A; covers team grants
Series A investors	~15.1%	Closer to 20% pre-ESOP-refresh; 15.1% post
Founders combined: ~47.0%	**100%**	**Right at the line where founder control begins to matter**

The single number to anchor on. If everything goes to plan, founders together hold ~47% post-Series-A. Each founder ~23.5%. That is **healthy and normal** for two-founder Indian/global SaaS — many founders end Series A at 15–20% each. The discipline to preserve this is: **don't raise more than you need at each round** (which is why $1.2–1.5M, not $3M, at angel). Every extra $500K raised at this stage costs ~3–4% in dilution that compounds across rounds.

PART 8 OF 8

08

Risks & sensitivities

What breaks the plan, by how much, and what we'd do.

8.1 · Top six risks, ranked

Risk	Likelihood	Impact	Mitigation
PoC→paid conversion is slower than 30%	Medium	High	Tighten qualification; require written success criteria; shorter PoCs. Build a second PoC offering for non-converted accounts.
Founding sales hire underperforms (long ramp)	Medium-high	High	Founder-led selling for first 10 customers; hire sales only after PMF signal. Don't hire sales to *find* PMF.
Top engineer cannot be hired at planned salary	Medium	Medium	Higher ESOP % for first 3 senior hires (above the 10% pool). Trade equity for cash.
AI gateway costs spike if usage scales fast	Low-medium	Medium	Accelerate the fine-tune timeline; route low-stakes traffic to self-hosted models earlier.
Competitor (Veeva, MasterControl, ComplianceQuest) launches India-specific tier	Low	Medium-high	Speed matters. Lock 10+ reference customers before they react. Differentiated on AI + supplier coupling, not just price.
Regulatory shifts (e.g. data localization) change cost base	Low	Variable	On-prem deployment is already part of the roadmap; we can support hosting requirements without architecture change.

8.2 · Sensitivity analysis — ARR at M18 if assumptions move

Variable	Base	Pessimistic	Optimistic	M18 ARR delta
Customers at M18	30	18	45	-$100K to +$130K
Blended ACV	$8.5K	$6.5K	$11K	-$60K to +$75K
PoC→paid conversion	35%	22%	50%	Drives customer count above
Sales cycle (months)	4	7	2.5	Shifts ARR right or left

Even in the pessimistic case, M18 ARR is ~$120–150K — enough to make a seed-round case, just at lower valuation. The plan does not require optimistic assumptions to work; it requires base-case execution.

8.3 · What we'd do at M9 if behind plan

1. Freeze the M9 hire (sales) until pipeline is real.
2. Founder-led sales sprint — 50 outbound calls/week, 10 demos/week for 8 weeks.
3. Pivot positioning: lead with the audit-prep agent only (one wedge), not full EQMS. Easier sale.
4. Re-focus on Tier 3 CDMOs (highest pain, fastest close) until 5 paid customers achieved.
5. If still nothing at M12: serious conversation with investors. Bridge round, pivot, or wind-down preparation. Don't fundraise on a broken plan.

> **Closing thought.** The plan above is the **conservative** case. The optimistic case (~$2.5–3M ARR at M36, 200+ customers) doesn't change the funding ask — it changes the seed/Series-A timing and terms. The downside case (~$0.5M ARR, 60 customers at M36) means a smaller Series A or extension round, but the unit economics still work and the company survives. The mid-case is what we're underwriting and what the $1.2–1.5M angel funds. **This is what the deck's $3M ask should have been.**